

Bigmart_new

June 20, 2022

```
[334]: #importing pandas liabries
import pandas as pd
import numpy as np
import seaborn as sns
import statsmodels.api as sm
import statsmodels.formula.api as smf
from sklearn import linear_model
import warnings
warnings.filterwarnings('ignore')
```

```
[335]: #importing the data
data_test=pd.read_csv('bigmart_test.csv')
data_train=pd.read_csv('bigmart_train.csv')
```

```
[336]: data_train
```

```
[336]:
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	\
0	FDA15	9.300	Low Fat	0.016047	
1	DRC01	5.920	Regular	0.019278	
2	FDN15	17.500	Low Fat	0.016760	
3	FDX07	19.200	Regular	0.000000	
4	NCD19	8.930	Low Fat	0.000000	
...	
8518	FDF22	6.865	Low Fat	0.056783	
8519	FDS36	8.380	Regular	0.046982	
8520	NCJ29	10.600	Low Fat	0.035186	
8521	FDN46	7.210	Regular	0.145221	
8522	DRG01	14.800	Low Fat	0.044878	

	Item_Type	Item_MRP	Outlet_Identifier	\
0	Dairy	249.8092	OUT049	
1	Soft Drinks	48.2692	OUT018	
2	Meat	141.6180	OUT049	
3	Fruits and Vegetables	182.0950	OUT010	
4	Household	53.8614	OUT013	
...	
8518	Snack Foods	214.5218	OUT013	

8519	Baking Goods	108.1570	OUT045
8520	Health and Hygiene	85.1224	OUT035
8521	Snack Foods	103.1332	OUT018
8522	Soft Drinks	75.4670	OUT046

	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	\
0	1999	Medium	Tier 1	
1	2009	Medium	Tier 3	
2	1999	Medium	Tier 1	
3	1998	NaN	Tier 3	
4	1987	High	Tier 3	
...	
8518	1987	High	Tier 3	
8519	2002	NaN	Tier 2	
8520	2004	Small	Tier 2	
8521	2009	Medium	Tier 3	
8522	1997	Small	Tier 1	

	Outlet_Type	Item_Outlet_Sales
0	Supermarket Type1	3735.1380
1	Supermarket Type2	443.4228
2	Supermarket Type1	2097.2700
3	Grocery Store	732.3800
4	Supermarket Type1	994.7052
...
8518	Supermarket Type1	2778.3834
8519	Supermarket Type1	549.2850
8520	Supermarket Type1	1193.1136
8521	Supermarket Type2	1845.5976
8522	Supermarket Type1	765.6700

[8523 rows x 12 columns]

```
[337]: #Summary statistics of numerical data
data_train.describe()
```

```
[337]:
```

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	\
count	7060.000000	8523.000000	8523.000000	8523.000000	
mean	12.857645	0.066132	140.992782	1997.831867	
std	4.643456	0.051598	62.275067	8.371760	
min	4.555000	0.000000	31.290000	1985.000000	
25%	8.773750	0.026989	93.826500	1987.000000	
50%	12.600000	0.053931	143.012800	1999.000000	
75%	16.850000	0.094585	185.643700	2004.000000	
max	21.350000	0.328391	266.888400	2009.000000	

Item_Outlet_Sales

```

count      8523.000000
mean       2181.288914
std        1706.499616
min         33.290000
25%        834.247400
50%        1794.331000
75%        3101.296400
max        13086.964800

```

```
[338]: #Summary statistics for categorical data
data_train.describe(include=object).T
```

```
[338]:
```

	count	unique	top	freq
Item_Identifier	8523	1559	FDW13	10
Item_Fat_Content	8523	5	Low Fat	5089
Item_Type	8523	16	Fruits and Vegetables	1232
Outlet_Identifier	8523	10	OUT027	935
Outlet_Size	6113	3	Medium	2793
Outlet_Location_Type	8523	3	Tier 3	3350
Outlet_Type	8523	4	Supermarket Type1	5577

```
[339]: #printing the shape and of the train and test data
print(data_train.shape)
print(data_test.shape)
```

```

(8523, 12)
(5681, 11)

```

```
[340]: #Cols that not required for any analysis
data_train=data_train.drop(['Outlet_Identifier','Item_Identifier'],axis=1)
```

```
[341]: #Check and drop deplicate columns
print('Before dedup:', data_train.shape)
data_cln=data_train.loc[:, ~data_train.columns.duplicated()]
print('After dedup:', data_cln.shape)
duplicateCols=data_train.loc[:, data_train.columns.duplicated()]
if (duplicateCols.shape[1] !=0):
    print('Number of duplicated columns dropped:', duplicatedCols.shape[1])
    print("Dupliate columns except first occurrences:")
    print(list(duplicateCols.columns))
```

```

Before dedup: (8523, 10)
After dedup: (8523, 10)

```

```
[342]: #Check and drop duplicated rows based on all columns
print ('Before dedup:', data_cln.shape)
data_cln.drop_duplicates(inplace=True) #by default keep='first'
```

```

print('After dedup:', data_cln.shape)
duplicateRows=data_train[data_train.duplicated()]
print('Number of duplicated rows dropped:', data_train.shape[0] - data_cln.
↳shape[0])
if (data_cln.shape[0] - data_train.shape[0] !=0):
    print("Duplicate rows except first occurrence:")
    print(duplicateRows)

```

Before dedup: (8523, 10)

After dedup: (8523, 10)

Number of duplicated rows dropped: 0

```

[343]: #Finding Missing values
def zero_missing_unique(df):
    var_miss_rate=df.isnull().sum(axis=0)/df.shape[0]
    var_miss_rate=var_miss_rate.to_frame('missing_rate')

    var_zero_rate=(df==0).astype(int).sum(axis=0)/df.shape[0]
    var_zero_rate=var_zero_rate.to_frame('zero_rate')
    var_unique=df.nunique().to_frame('count_unique')
    var_type=df.dtypes.to_frame('data_type')
    data_stat=pd.concat([var_miss_rate, var_zero_rate,var_unique, var_type],
↳axis=1)
    return data_stat
stat_zero_missing_unique=zero_missing_unique(data_train)# just change the data_
↳name here to your own data
stat_zero_missing_unique

```

```

[343]:

```

	missing_rate	zero_rate	count_unique	data_type
Item_Weight	0.171653	0.000000	415	float64
Item_Fat_Content	0.000000	0.000000	5	object
Item_Visibility	0.000000	0.061715	7880	float64
Item_Type	0.000000	0.000000	16	object
Item_MRP	0.000000	0.000000	5938	float64
Outlet_Establishment_Year	0.000000	0.000000	9	int64
Outlet_Size	0.282764	0.000000	3	object
Outlet_Location_Type	0.000000	0.000000	3	object
Outlet_Type	0.000000	0.000000	4	object
Item_Outlet_Sales	0.000000	0.000000	3493	float64

```

[344]: #Dealing with the missing rate
#do not drop high missing rate variable
#use imputing approach specified in preprocessing
stat_zero_missing_unique.sort_values(by='missing_rate', ascending=False)

```

```

[344]:

```

	missing_rate	zero_rate	count_unique	data_type
Outlet_Size	0.282764	0.000000	3	object

Item_Weight	0.171653	0.000000	415	float64
Item_Fat_Content	0.000000	0.000000	5	object
Item_Visibility	0.000000	0.061715	7880	float64
Item_Type	0.000000	0.000000	16	object
Item_MRP	0.000000	0.000000	5938	float64
Outlet_Establishment_Year	0.000000	0.000000	9	int64
Outlet_Location_Type	0.000000	0.000000	3	object
Outlet_Type	0.000000	0.000000	4	object
Item_Outlet_Sales	0.000000	0.000000	3493	float64

```
[345]: #Grouping by Item Type
item_type_baseline = data_train.groupby('Item_Type').mean()[['Item_Weight']].
    ↪reset_index()
item_type_baseline = item_type_baseline.to_numpy()
item_type_baseline
```

```
[345]: array([[ 'Baking Goods', 12.277108208955255],
      [ 'Breads', 11.3469362745098],
      [ 'Breakfast', 12.768202247191002],
      [ 'Canned', 12.305705009276451],
      [ 'Dairy', 13.42606890459367],
      [ 'Frozen Foods', 12.867061281337076],
      [ 'Fruits and Vegetables', 13.224769381746881],
      [ 'Hard Drinks', 11.400327868852452],
      [ 'Health and Hygiene', 13.142313953488392],
      [ 'Household', 13.384736495388726],
      [ 'Meat', 12.81734421364986],
      [ 'Others', 13.853284671532842],
      [ 'Seafood', 12.5528431372549],
      [ 'Snack Foods', 12.987879554655919],
      [ 'Soft Drinks', 11.847459893048129],
      [ 'Starchy Foods', 13.690730769230765]], dtype=object)
```

```
[346]: ## Item weight standardization using item type
for x in item_type_baseline:
    print(x[0])
```

Baking Goods
 Breads
 Breakfast
 Canned
 Dairy
 Frozen Foods
 Fruits and Vegetables
 Hard Drinks
 Health and Hygiene
 Household

Meat
Others
Seafood
Snack Foods
Soft Drinks
Starchy Foods

```
[347]: def item_weight_cal(cols):  
        Item_Wt = cols[0]  
        Item_Tp = cols[1]  
        if pd.isnull(Item_Wt):  
            for x in item_type_baseline:  
                if (x[0] == Item_Tp):  
                    # print (x[1])  
                    return x[1]  
        else:  
            return Item_Wt
```

```
[348]: data_train['Item_Weight'] =data_train[['Item_Weight','Item_Type']].  
        ↪apply(item_weight_cal, axis = 1)  
data_train.isna().sum()
```

```
[348]: Item_Weight                0  
Item_Fat_Content                0  
Item_Visibility                0  
Item_Type                      0  
Item_MRP                       0  
Outlet_Establishment_Year      0  
Outlet_Size                    2410  
Outlet_Location_Type           0  
Outlet_Type                    0  
Item_Outlet_Sales              0  
dtype: int64
```

```
[349]: data_train[['Outlet_Type','Outlet_Size']].  
        ↪groupby(['Outlet_Type','Outlet_Size']).size()
```

```
[349]: Outlet_Type    Outlet_Size  
Grocery Store    Small          528  
Supermarket Type1 High          932  
                  Medium         930  
                  Small        1860  
Supermarket Type2 Medium         928  
Supermarket Type3 Medium         935  
dtype: int64
```

```
[350]: df = data_train[['Outlet_Type', 'Outlet_Size']]
df.drop_duplicates()
```

```
[350]:      Outlet_Type Outlet_Size
0   Supermarket Type1      Medium
1   Supermarket Type2      Medium
3      Grocery Store         NaN
4   Supermarket Type1      High
7   Supermarket Type3      Medium
8   Supermarket Type1         NaN
11  Supermarket Type1      Small
23      Grocery Store      Small
```

```
[351]: #Replacing the Missing values with small in the data
data_train['Outlet_Size']=data_train['Outlet_Size'].fillna('Small')
```

```
[352]: data_train.isna().sum()
```

```
[352]: Item_Weight      0
Item_Fat_Content    0
Item_Visibility     0
Item_Type           0
Item_MRP            0
Outlet_Establishment_Year  0
Outlet_Size         0
Outlet_Location_Type  0
Outlet_Type         0
Item_Outlet_Sales    0
dtype: int64
```

```
[353]: data_train.groupby(['Item_Fat_Content']).size()
```

```
[353]: Item_Fat_Content
LF      316
Low Fat 5089
Regular 2889
low fat  112
reg      117
dtype: int64
```

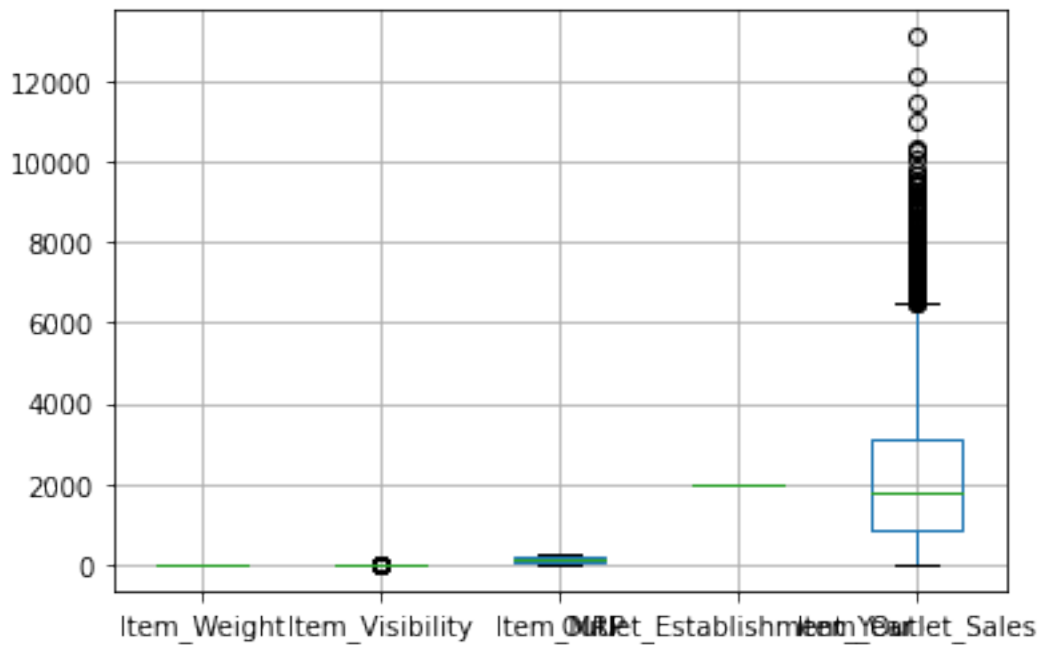
```
[354]: #Observed irregular data pattern where low fat is represented as LF and low fat
→and all regular presented as reg
data_train["Item_Fat_Content"].replace({"LF": "Low Fat", "low fat": "Low_Fat",
→"reg": "Regular"}, inplace=True)
```

```
[355]: data_train.groupby(['Item_Fat_Content']).size()
```

```
[355]: Item_Fat_Content
Low Fat    5405
Low_Fat    112
Regular    3006
dtype: int64
```


```
[356]: #Checking for Outlier using box plot for the entire data
data_train.boxplot()
```

```
[356]: <AxesSubplot:>
```



```
[357]: sns.boxplot(data_train['Item_Visibility'])
```

```
[357]: <AxesSubplot:xlabel='Item_Visibility'>
```


```
[358]: #Putting a cap on the data
data_train[data_train['Item_Visibility']<0.175].count()
```

```
[358]: Item_Weight      8245
Item_Fat_Content    8245
Item_Visibility     8245
Item_Type           8245
Item_MRP            8245
Outlet_Establishment_Year 8245
Outlet_Size         8245
Outlet_Location_Type 8245
Outlet_Type         8245
Item_Outlet_Sales   8245
dtype: int64
```

```
[359]: #Viewing the Cap
data_train= data_train[data_train['Item_Visibility'].values < 0.175]
```


```
[360]: #Viewing the boxplot after the data treatment
sns.boxplot(data_train['Item_Visibility'])
```

```
[360]: <AxesSubplot:xlabel='Item_Visibility'>
```




```
[361]: #Viewing the box plot for Weight  
sns.boxplot(data_train['Item_Weight'])
```

```
[361]: <AxesSubplot:xlabel='Item_Weight'>
```




```
[362]: sns.boxplot(data_train['Item_MRP'])
```

```
[362]: <AxesSubplot:xlabel='Item_MRP'>
```



```
[363]: #viewing box plot for item outlet sales  
sns.boxplot(data_train['Item_Outlet_Sales'])
```

```
[363]: <AxesSubplot:xlabel='Item_Outlet_Sales'>
```




```
[364]: #Placing a Cap on the data
data_train[data_train['Item_Outlet_Sales'].values<6000].count()
```

```
[364]: Item_Weight          7958
Item_Fat_Content        7958
Item_Visibility         7958
Item_Type               7958
Item_MRP                7958
Outlet_Establishment_Year 7958
Outlet_Size            7958
Outlet_Location_Type    7958
Outlet_Type            7958
Item_Outlet_Sales       7958
dtype: int64
```

```
[365]: #Viewing the data after the Cap
data_train=data_train[data_train['Item_Outlet_Sales'].values<6000]
```

```
[366]: #Viewing the box plot after the treatment
sns.boxplot(data_train['Item_Outlet_Sales'])
```

```
[366]: <AxesSubplot:xlabel='Item_Outlet_Sales'>
```



```
[367]: #check variable with extreme high correlation with target variable for
        ↳potential target leakage
print("correlation with target variable:\n")
data_train.corr()['Item_Outlet_Sales'].sort_values(ascending=False)
```

correlation with target variable:

```
[367]: Item_Outlet_Sales      1.000000
       Item_MRP              0.554841
       Item_Weight           0.009989
       Outlet_Establishment_Year -0.001486
       Item_Visibility        -0.063880
       Name: Item_Outlet_Sales, dtype: float64
```

```
[368]: data_corr=data_train.corr()
       data_corr
```


```
[368]:
```

	Item_Weight	Item_Visibility	Item_MRP	\
Item_Weight	1.000000	-0.005938	0.024639	
Item_Visibility	-0.005938	1.000000	0.002651	
Item_MRP	0.024639	0.002651	1.000000	
Outlet_Establishment_Year	-0.010047	-0.032005	0.029837	
Item_Outlet_Sales	0.009989	-0.063880	0.554841	

	Outlet_Establishment_Year	Item_Outlet_Sales
Item_Weight	-0.010047	0.009989
Item_Visibility	-0.032005	-0.063880
Item_MRP	0.029837	0.554841
Outlet_Establishment_Year	1.000000	-0.001486
Item_Outlet_Sales	-0.001486	1.000000

```
[369]: #correlation heatmap
import matplotlib as plt
%matplotlib inline
sns.heatmap(data =data_corr ,square = True)
```

[369]: <AxesSubplot:>



```
[370]: from sklearn.preprocessing import MinMaxScaler, StandardScaler, LabelEncoder, \
        Normalizer, RobustScaler
```

```
[371]: #Grouping the
cols = ['Item_Fat_Content', 'Outlet_Size', 'Outlet_Type',
        'Outlet_Location_Type', 'Outlet_Establishment_Year', 'Item_Type']
```

```
[372]: #Applying the label encoder
data_train[cols] = data_train[cols].apply(LabelEncoder().fit_transform)
```

```
[373]: data_train.head()
```

```
[373]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	\
0	9.30	0	0.016047	4	249.8092	
1	5.92	2	0.019278	14	48.2692	
2	17.50	0	0.016760	10	141.6180	
3	19.20	2	0.000000	6	182.0950	
4	8.93	0	0.000000	9	53.8614	

	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type	\
0	4	1	0	1	
1	8	1	2	2	
2	4	1	0	1	
3	3	2	2	0	
4	1	0	2	1	

	Item_Outlet_Sales
0	3735.1380
1	443.4228
2	2097.2700
3	732.3800
4	994.7052

```
[374]: # Applying Min Max Scaler
min_max_scal = MinMaxScaler()
min_max_scal.fit(data_train[['Item_Outlet_Sales']])
min_max_scal.data_max_
data_train['Item_Outlet_Sales_Max'] = min_max_scal.
        transform(data_train[['Item_Outlet_Sales']])
data_train.head()
```

```
[374]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	\
0	9.30	0	0.016047	4	249.8092	
1	5.92	2	0.019278	14	48.2692	
2	17.50	0	0.016760	10	141.6180	
3	19.20	2	0.000000	6	182.0950	
4	8.93	0	0.000000	9	53.8614	

	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type	\
0	4	1	0	1	

1	8	1	2	2
2	4	1	0	1
3	3	2	2	0
4	1	0	2	1

	Item_Outlet_Sales	Item_Outlet_Sales_Max
0	3735.1380	0.621229
1	443.4228	0.068827
2	2097.2700	0.346369
3	732.3800	0.117318
4	994.7052	0.161341

```
[375]: ## Applying Standard Scaler
std_scal = StandardScaler()
std_scal.fit(data_train[['Item_Outlet_Sales']])
std_scal.mean_
data_train['Item_Outlet_Sales_Std'] = std_scal.
↳transform(data_train[['Item_Outlet_Sales']])
data_train.head()
```

```
[375]: Item_Weight  Item_Fat_Content  Item_Visibility  Item_Type  Item_MRP  \
0          9.30             0          0.016047         4  249.8092
1          5.92             2          0.019278        14   48.2692
2         17.50             0          0.016760        10  141.6180
3         19.20             2          0.000000         6  182.0950
4          8.93             0          0.000000         9   53.8614
```

	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type	\
0	4	1	0	1	
1	8	1	2	2	
2	4	1	0	1	
3	3	2	2	0	
4	1	0	2	1	

	Item_Outlet_Sales	Item_Outlet_Sales_Max	Item_Outlet_Sales_Std
0	3735.1380	0.621229	1.177460
1	443.4228	0.068827	-1.116200
2	2097.2700	0.346369	0.036197
3	732.3800	0.117318	-0.914856
4	994.7052	0.161341	-0.732068

```
[376]: # Applying Normalizer Scaler
norm_scal = Normalizer()
norm_scal.fit(data_train[['Item_Outlet_Sales']])
data_train['Item_Outlet_Sales_Norm'] = norm_scal.
↳transform(data_train[['Item_Outlet_Sales']])
data_train.head()
```



```
[376]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	\
0	9.30	0	0.016047	4	249.8092	
1	5.92	2	0.019278	14	48.2692	
2	17.50	0	0.016760	10	141.6180	
3	19.20	2	0.000000	6	182.0950	
4	8.93	0	0.000000	9	53.8614	

	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type	\
0	4	1	0	1	
1	8	1	2	2	
2	4	1	0	1	
3	3	2	2	0	
4	1	0	2	1	

	Item_Outlet_Sales	Item_Outlet_Sales_Max	Item_Outlet_Sales_Std	\
0	3735.1380	0.621229	1.177460	
1	443.4228	0.068827	-1.116200	
2	2097.2700	0.346369	0.036197	
3	732.3800	0.117318	-0.914856	
4	994.7052	0.161341	-0.732068	

	Item_Outlet_Sales_Norm
0	1.0
1	1.0
2	1.0
3	1.0
4	1.0

```
[377]: # Applying Robust Scaler
rob_scal = RobustScaler()
rob_scal.fit(data_train[['Item_Outlet_Sales']])
rob_scal.scale_
data_train['Item_Outlet_Sales_Rob'] = rob_scal.
↳transform(data_train[['Item_Outlet_Sales']])
data_train
```

```
[377]:
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	\
0	9.300	0	0.016047	4	249.8092	
1	5.920	2	0.019278	14	48.2692	
2	17.500	0	0.016760	10	141.6180	
3	19.200	2	0.000000	6	182.0950	
4	8.930	0	0.000000	9	53.8614	
...	
8518	6.865	0	0.056783	13	214.5218	
8519	8.380	2	0.046982	0	108.1570	
8520	10.600	0	0.035186	8	85.1224	
8521	7.210	2	0.145221	13	103.1332	

8522	14.800	0	0.044878	14	75.4670
------	--------	---	----------	----	---------

	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	\
0	4	1	0	
1	8	1	2	
2	4	1	0	
3	3	2	2	
4	1	0	2	
...	
8518	1	0	2	
8519	5	2	1	
8520	6	2	1	
8521	8	1	2	
8522	2	2	0	

	Outlet_Type	Item_Outlet_Sales	Item_Outlet_Sales_Max	\
0	1	3735.1380	0.621229	
1	2	443.4228	0.068827	
2	1	2097.2700	0.346369	
3	0	732.3800	0.117318	
4	1	994.7052	0.161341	
...	
8518	1	2778.3834	0.460670	
8519	1	549.2850	0.086592	
8520	1	1193.1136	0.194637	
8521	2	1845.5976	0.304134	
8522	1	765.6700	0.122905	

	Item_Outlet_Sales_Std	Item_Outlet_Sales_Norm	Item_Outlet_Sales_Rob
0	1.177460	1.0	0.930774
1	-1.116200	1.0	-0.624921
2	0.036197	1.0	0.156702
3	-0.914856	1.0	-0.488357
4	-0.732068	1.0	-0.364380
...
8518	0.510796	1.0	0.478603
8519	-1.042436	1.0	-0.574890
8520	-0.593817	1.0	-0.270610
8521	-0.139168	1.0	0.037760
8522	-0.891659	1.0	-0.472624

[7958 rows x 14 columns]

```
[378]: #Assiging Features and Target variables
#Using this alternative code print('data before splitting:', data_train.shape)
#x_cols=[x for x in data_train.columns if x!='Item_Outlet_Sales']
```

```

features = data_train.
↳drop(['Item_Outlet_Sales','Item_Outlet_Sales_Max','Item_Outlet_Sales_Std','Item_Outlet_Sale
target= data_train[['Item_Outlet_Sales_Std']]
features.head()

```

```

[378]:
Item_Weight  Item_Fat_Content  Item_Visibility  Item_Type  Item_MRP  \
0          9.30              0          0.016047          4  249.8092
1          5.92              2          0.019278         14   48.2692
2         17.50              0          0.016760         10  141.6180
3         19.20              2          0.000000          6  182.0950
4          8.93              0          0.000000          9   53.8614

Outlet_Establishment_Year  Outlet_Size  Outlet_Location_Type  Outlet_Type
0                        4              1              0              1
1                        8              1              2              2
2                        4              1              0              1
3                        3              2              2              0
4                        1              0              2              1

```

```

[379]: from sklearn.linear_model import Lasso, Ridge, ElasticNet, LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

```

```

[380]: X_train, X_test, y_train, y_test = train_test_split(features,
↳target, random_state = 10)

```

```

[381]: #Viewing the shape of the data
x_train.shape, x_test.shape, y_train.shape, y_test.shape

```

```

[381]: ((6366, 9), (1592, 9), (5968, 1), (1990, 1))

```

```

[382]: my_multiple_lm_simple = LinearRegression()
my_multiple_lm_simple.fit(X_train, y_train)
my_multiple_lm_simple.coef_
my_multiple_lm_simple.intercept_
my_multiple_lm_simple.preds = my_multiple_lm_simple.predict(X_test)
mean_absolute_error(my_multiple_lm_simple.preds, y_test)

```

```

[382]: 0.5649110474262632

```

```

[383]: #Finding the Mean Square Error
mean_squared_error(my_multiple_lm_simple.preds, y_test)

```

```

[383]: 0.5301009454480488

```

```

[384]: #getting the R_squared
r2_score(y_test, my_multiple_lm_simple.preds)

```

[384]: 0.46042229424616143

```
[385]: #Finding the Lasso
alphas=[-5,-1,1e-8,1e-6,1e-4,1e-3,1e-2,1,5]

def test_alpha(a):
    model_ridge = Ridge(alpha=a)
    model_ridge.fit(X_train, y_train)
    pred_test_ridge = model_ridge.predict(X_test)
    new_score = r2_score(y_test, pred_test_ridge)
    new_mse = mean_squared_error(y_test, pred_test_ridge)
    print('ALPHA: {:.3f} R2 SCORE: {:.4f}% new_score, {:.1f}'.
    ↪format(a,new_score, new_mse))

for alpha in alphas:
    test_alpha(alpha)
```

```
ALPHA: -5.000 R2 SCORE: 0.4606% new_score, 0.5
ALPHA: -1.000 R2 SCORE: 0.4605% new_score, 0.5
ALPHA: 0.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.001 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.010 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 1.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 5.000 R2 SCORE: 0.4601% new_score, 0.5
```

```
[386]: #Finding the Lasso
alphas =[-5, -1, 1e-8,1e-6,1e-4, 1e-3, 1e-2, 1, 5]

def test_alpha(a):
    model_lasso = Lasso(alpha=a)
    model_lasso.fit(X_train, y_train)
    pred_test_lasso = model_lasso.predict(X_test)
    new_score = r2_score(y_test, pred_test_lasso)
    new_mse = mean_squared_error(y_test, pred_test_lasso)
    print('ALPHA: {:.3f} R2 SCORE: {:.4f}% new_score, {:.1f}'.
    ↪format(a,new_score, new_mse))

for alpha in alphas:
    test_alpha(alpha)
```

```
ALPHA: -5.000 R2 SCORE: -14270.5003% new_score, 14020.8
ALPHA: -1.000 R2 SCORE: -569.7404% new_score, 560.7
ALPHA: 0.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.001 R2 SCORE: 0.4592% new_score, 0.5
```

```
ALPHA: 0.010 R2 SCORE: 0.4585% new_score, 0.5
ALPHA: 1.000 R2 SCORE: 0.3115% new_score, 0.7
ALPHA: 5.000 R2 SCORE: 0.3071% new_score, 0.7
```

```
[387]: #Finding the ElasticNet
alphas = [-5, -1, 1e-8, 1e-6, 1e-4, 1e-3, 1e-2, 1, 5]

def test_alpha(a):
    model_elastic = ElasticNet(alpha=a)
    model_elastic.fit(X_train, y_train)
    pred_test_elastic = model_elastic.predict(X_test)
    new_score = r2_score(y_test, pred_test_elastic)
    new_mse = mean_squared_error(y_test, pred_test_elastic)
    print('ALPHA: {:.3f} R2 SCORE: {:.4f}% new_score, {:.1f}'.
        ↪format(a, new_score, new_mse))

for alpha in alphas:
    test_alpha(alpha)
```

```
ALPHA: -5.000 R2 SCORE: -17.1553% new_score, 17.8
ALPHA: -1.000 R2 SCORE: -0.0005% new_score, 1.0
ALPHA: 0.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.000 R2 SCORE: 0.4604% new_score, 0.5
ALPHA: 0.001 R2 SCORE: 0.4598% new_score, 0.5
ALPHA: 0.010 R2 SCORE: 0.4590% new_score, 0.5
ALPHA: 1.000 R2 SCORE: 0.3114% new_score, 0.7
ALPHA: 5.000 R2 SCORE: 0.3110% new_score, 0.7
```

```
[388]: #Usng Manual Calculation
my_ols_model = smf.ols(formula='Item_Outlet_Sales_Std ~ Item_Weight_
    ↪+Item_Fat_Content + Item_Visibility + Item_Type + Item_MRP_
    ↪+Outlet_Establishment_Year + Outlet_Size + Outlet_Location_Type_
    ↪+Outlet_Type', data =data_train).fit()
my_ols_model.summary()
```

```
[388]: <class 'statsmodels.iolib.summary.Summary'>
      """
                                OLS Regression Results
=====
=
Dep. Variable:      Item_Outlet_Sales_Std   R-squared:
0.467
Model:                                OLS   Adj. R-squared:
0.466
Method:                                Least Squares   F-statistic:
773.3
```

Date: Sun, 27 Feb 2022 Prob (F-statistic):
0.00
Time: 19:32:39 Log-Likelihood:
-8789.4
No. Observations: 7958 AIC:
1.760e+04
Df Residuals: 7948 BIC:
1.767e+04
Df Model: 9
Covariance Type: nonrobust

		coef	std err	t	P> t
[0.025 0.975]					

Intercept		-1.5677	0.049	-31.953	0.000
-1.664	-1.472				
Item_Weight		-0.0003	0.002	-0.178	0.859
-0.004	0.003				
Item_Fat_Content		0.0215	0.009	2.470	0.014
0.004	0.039				
Item_Visibility		-0.6496	0.187	-3.466	0.001
-1.017	-0.282				
Item_Type		0.0013	0.002	0.642	0.521
-0.003	0.005				
Item_MRP		0.0093	0.000	69.528	0.000
0.009	0.010				
Outlet_Establishment_Year		0.0166	0.003	5.067	0.000
0.010	0.023				
Outlet_Size		-0.1228	0.015	-8.199	0.000
-0.152	-0.093				
Outlet_Location_Type		-0.2243	0.013	-17.712	0.000
-0.249	-0.200				
Outlet_Type		0.5498	0.012	44.180	0.000
0.525	0.574				
=====					
Omnibus:		109.554	Durbin-Watson:		2.030
Prob(Omnibus):		0.000	Jarque-Bera (JB):		117.495
Skew:		0.265	Prob(JB):		3.06e-26
Kurtosis:		3.273	Cond. No.		3.48e+03
=====					

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.48e+03. This might indicate that there are

```
strong multicollinearity or other numerical problems.
"""
```

```
[389]: #Polynomial Linear Regression
```

```
from sklearn.preprocessing import PolynomialFeatures
features_poly=PolynomialFeatures(degree=2)
train_Poly=features_poly.fit_transform(features)
train_Poly
```

```
[389]: array([[ 1. ,  9.3 ,  0. , ...,  0. ,  0. ,  1. ],
 [ 1. ,  5.92,  2. , ...,  4. ,  4. ,  4. ],
 [ 1. , 17.5 ,  0. , ...,  0. ,  0. ,  1. ],
 ...,
 [ 1. , 10.6 ,  0. , ...,  1. ,  1. ,  1. ],
 [ 1. ,  7.21,  2. , ...,  4. ,  4. ,  4. ],
 [ 1. , 14.8 ,  0. , ...,  0. ,  0. ,  1. ]])
```

```
[390]: #Splitting the data
```

```
X_train, X_test, y_train, y_test = train_test_split(train_Poly,
↳target,random_state = 6)
```

```
[391]: #Fitting the Regression
```

```
my_model = LinearRegression()
my_model.fit(X_train, y_train)
preds = my_model.predict(X_test)
print('MAE : ', mean_absolute_error(y_test, preds))
print('MSE : ', mean_squared_error(y_test, preds))
print('R2 : ', r2_score(y_test, preds))
```

```
MAE :  0.5063635221381928
MSE :  0.44709466714578755
R2 :  0.5517740971291534
```

```
[392]: #Printing the shape of the data
```

```
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(5968, 55)
(1990, 55)
(5968, 1)
(1990, 1)
```

```
[398]: from sklearn.decomposition import PCA
```

```
[399]: #using the PCA  
pca=PCA(n_components=30)  
pca.fit(X_train)
```

```
[399]: PCA(n_components=30)
```

```
[400]: #Finding the explained variance  
pca.explained_variance_ratio_
```

```
[400]: array([9.96552934e-01, 1.51322829e-03, 1.20511587e-03, 5.27418264e-04,  
7.62093643e-05, 6.37594963e-05, 2.42569520e-05, 1.89479857e-05,  
1.00901022e-05, 4.15432583e-06, 1.35881615e-06, 9.12735851e-07,  
3.92834018e-07, 3.38921562e-07, 1.78563094e-07, 1.49042067e-07,  
1.33364324e-07, 1.25299539e-07, 5.98842032e-08, 5.24073170e-08,  
4.95592705e-08, 4.19868843e-08, 2.33278668e-08, 1.95851504e-08,  
1.71995105e-08, 1.63838468e-08, 3.43722850e-09, 3.38018148e-09,  
1.74664954e-09, 1.57487884e-09])
```

```
[ ]:
```